

Sample Document

Neo4all Test Fixture | Version 1.0

Background

Knowledge graphs represent information as a network of entities and relationships, enabling machines to reason about complex domains. In a graph-based knowledge representation, nodes correspond to real-world objects such as organisations, people, locations, events, and concepts, while edges encode the semantic relationships that connect them. This structure supports a wide range of analytical tasks including entity resolution, link prediction, and question answering. Constructing a knowledge graph from unstructured text requires a pipeline that extracts named entities, resolves coreferences, and identifies relationships across sentences and documents. The quality of the resulting graph depends on the accuracy of each stage and on the schema that defines which entity and relationship types are considered relevant to the target domain.

This fixture document is used to verify that the ingestion pipeline correctly parses document content, extracts text with page locators, and produces deterministic chunk identifiers from the page-locator hash function.

Sample Table

Table 1 — Entity types extracted from knowledge graph literature

Entity Type	Primary Property	Example Values
Organisation	name	Acme Corp, OpenAI, WHO
Person	full_name	Jane Smith, John Doe
Location	place_name	London, New York, Geneva
Event	title	Annual Conference 2024
Concept	label	Machine Learning, NLP

Source: neo4all fixtures — for testing only